

project

January 9, 2025

```
[147]: library(tidyr)
library("pxweb")
library("tidyverse")
library("GGally")
library("ggpubr")
library(gridExtra)
library(cowplot)
```

We are going to explore if it is reasonable to use a MLR model to predict the number of people graduating from three different levels of education. We will also explore if the model performs better for the entire population or for men and women separately.

```
[148]: # d <- pxweb_interactive("https://api.scb.se/OV0104/v1/doris/sv/ssd/UF/UF0550/
      ↪UF0550C/Historisk11bN")
API_wait <- 1e-3
pxq_edu <-
  list(
    "Examen" = c("*"), # Use "*" to select all
    "Kon" = c("*"),
    "ContentsCode" = c("000004NE"),
    "Tid" = c("*")
  )
pxg_edu <- pxweb_get("https://api.scb.se/OV0104/v1/doris/sv/ssd/START/UF/UF0550/
      ↪UF0550C/Historisk11bN", pxq_edu)
Sys.sleep(API_wait)
pxq_pop <-
  list(
    "Kon" = c("*"),
    "ContentsCode" = c("000000LV"),
    "Tid" = c("*")
  )
pxg_pop <- pxweb_get("https://api.scb.se/OV0104/v1/doris/sv/ssd/START/BE/BE0101/
      ↪BE0101G/BefUtvKon1749", pxq_pop)
Sys.sleep(API_wait)
# https://api.scb.se/OV0104/v1/doris/sv/ssd/START/HE/HE0103/HE0103A/ArbInk28
pxq_income <- list(
  "Kon" = c("*"),
```

```

  "Tid" = c("*"),
  "ContentsCode" = c("HE0103CL")
  # "HE0103CM")
)
pxg_income <- pxweb_get("https://api.scb.se/OV0104/v1/doris/sv/ssd/START/HE/
↪HE0103/HE0103A/ArbInk28", pxq_income)

Sys.sleep(API_wait)
pxq_kpi <- list(
  "Tid" = c("*"),
  "ContentsCode" = c("000004VU")
)
pxg_kpi <- pxweb_get("https://api.scb.se/OV0104/v1/doris/sv/ssd/START/PR/PR0101/
↪PR0101A/KPItotM", pxq_kpi)

```

We start by gathering data from SCB about education, income, KPI and population. Then we clean and modify the data to fit our needs.

```

[ ]: edu_df <- as.data.frame(pxg_edu)
edu_df <- as_tibble(edu_df)

doktor_df <- edu_df %>%
  filter(examen %in% c("Doktorsgrad", "Doktorsexamen")) %>%
  group_by(läsår, kön) %>%
  mutate(`Utfärdade examina vid universitet och högskolor` = sum(`Utfärdade
↪examina vid universitet och högskolor`)) %>%
  filter(examen == "Doktorsexamen") %>%
  ungroup()

licenciat_df <- edu_df %>%
  filter(examen %in% c("Licentiatexamen", "Licentiatexamen (Äldre)")) %>%
  group_by(läsår, kön) %>%
  mutate(`Utfärdade examina vid universitet och högskolor` = sum(`Utfärdade
↪examina vid universitet och högskolor`)) %>%
  filter(examen == "Licentiatexamen") %>%
  ungroup()

edu_df <- edu_df %>%
  filter(examen == "Examen från grundutbildning") %>%
  full_join(licenciat_df) %>%
  full_join(doktor_df) %>%
  pivot_wider(names_from = kön, values_from = `Utfärdade examina vid
↪universitet och högskolor`) %>%
  mutate(läsår = as.integer(str_extract(läsår, "^\\w*"))) %>%
  rename(år = läsår) %>%

```

```

pivot_longer(!c(år, examen), names_to = "kön", values_to = "antal") %>%
mutate(
  kön = str_replace(kön, "båda könen", "totalt"),
  examen = factor(examen,
    ordered = FALSE,
    levels = c(
      "Examen från grundutbildning", "Licentiatexamen",
      ↪ "Licentiatexamen (Äldre)", "Doktorsgrad", "Doktorsexamen"
    ),
    labels = c(
      "Grundutbildning", "Licentiatexamen", "Licentiatexamen",
      ↪ "Doktorsexamen", "Doktorsexamen"
    )
  )
)

pop_df <- as.data.frame(pwg_pop)
pop_df <- as_tibble(pop_df) %>% mutate(år = as.integer(år)) # %>%

income_df <- as.data.frame(pwg_income)
income_df <- as_tibble(income_df) %>%
  filter(år >= 1936) %>%
  mutate(år = as.integer(år))

kpi_df <- as.data.frame(pwg_kpi)
kpi_df <- as_tibble(kpi_df) %>%
  mutate(månad = str_extract(månad, "(20|19)[\\d]{2}")) %>%
  rename("år" = "månad") %>%
  group_by(år) %>%
  mutate(år = as.integer(år)) %>%
  summarize(`AVG_KPI_Fastställd` = mean(`KPI, fastställda tal`)) %>%
  ungroup()

final_df <- edu_df %>%
  left_join(income_df, by = join_by("år", kön == kön)) %>%
  left_join(kpi_df, by = "år") %>%
  left_join(pop_df, by = join_by("år", kön == kön)) %>%
  rename(medelinkomst = `Medelvärde, tkr`) %>%
  mutate(kön = factor(kön, levels = c("totalt", "kvinnor", "män"), ordered =
  ↪ FALSE)) %>%
  drop_na()

final_df %>% filter(år == 1980)

```

Joining with `by = join_by(examen, kön, läsår, `Utfärdade examina vid universitet och högskolor`)`

Joining with `by = join_by(examen, kön, läsår, `Utfärdade examina vid universitet och högskolor`)`

	examen	år	kön	antal	medelinkomst	AVG_KPI_Fastställd	Folkmän
	<fct>	<int>	<fct>	<dbl>	<dbl>	<dbl>	<dbl>
A tibble: 9 × 7	Grundutbildning	1980	totalt	34536	172.3	100	8317937
	Grundutbildning	1980	män	12110	220.4	100	4119822
	Grundutbildning	1980	kvinnor	22426	122.9	100	4198115
	Licentiatexamen	1980	totalt	1	172.3	100	8317937
	Licentiatexamen	1980	män	1	220.4	100	4119822
	Licentiatexamen	1980	kvinnor	0	122.9	100	4198115
	Doktorsexamen	1980	totalt	812	172.3	100	8317937
	Doktorsexamen	1980	män	666	220.4	100	4119822
	Doktorsexamen	1980	kvinnor	146	122.9	100	4198115

We then do some eploratory data analysis on the collected data

```
[150]: dok_df <- final_df %>%
  filter(examen == "Doktorsexamen", kön != "totalt") %>%
  mutate(Folkmängd = 1e-6 * Folkmängd)

lic_df <- final_df %>%
  filter(examen == "Licentiatexamen", kön != "totalt") %>%
  mutate(Folkmängd = 1e-6 * Folkmängd)

grund_df <- final_df %>%
  filter(examen == "Grundutbildning", kön != "totalt") %>%
  mutate(Folkmängd = 1e-6 * Folkmängd)

dok_pairs <- ggpairs(
  data = dok_df,
  columns = c("år", "antal", "medelinkomst", "AVG_KPI_Fastställd",
  ↪ "Folkmängd"),
  mapping = aes(color = kön)
)

lic_pairs <- ggpairs(
  data = lic_df, columns = c("år", "antal", "medelinkomst",
  ↪ "AVG_KPI_Fastställd", "Folkmängd"),
  mapping = aes(color = kön)
)

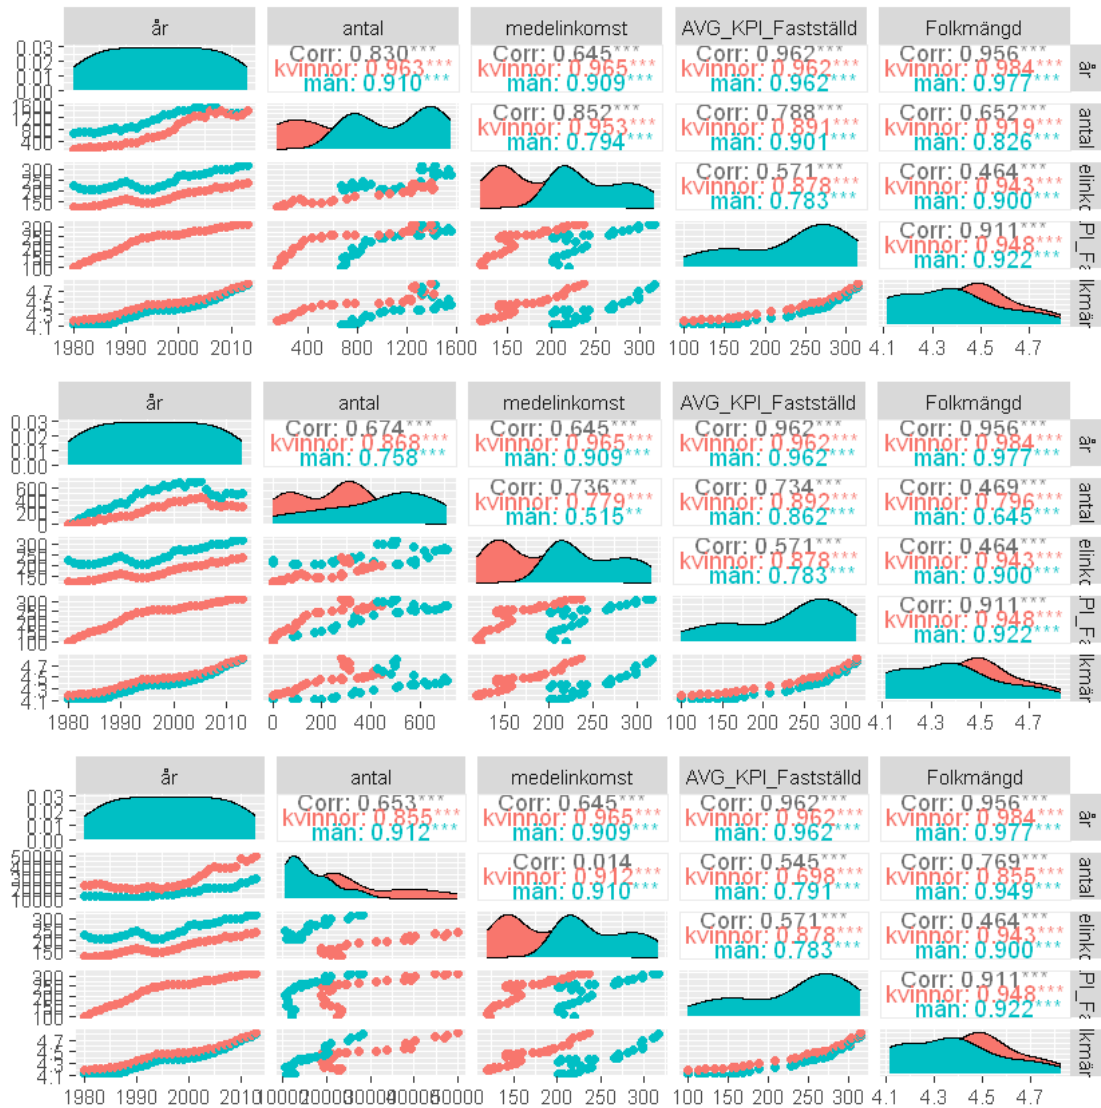
grund_pairs <- ggpairs(
  data = grund_df, columns = c("år", "antal", "medelinkomst",
  ↪ "AVG_KPI_Fastställd", "Folkmängd"),
  mapping = aes(color = kön)
```

```

)

plot_grid(
  ggmatrix_gtable(dok_pairs), ggmatrix_gtable(lic_pairs),
  ggmatrix_gtable(grund_pairs),
  ncol = 1
)

```



We can see that there is some collinearity between some of the columns and that it differs between different types of degrees.

As a naive approach the first model will contain all available data in the multiple linear regression for the amount of graduates.

```
[151]: naive_model <- lm(antal ~ år + kön + examen + medelinkomst + Folkmängd +
  ↪AVG_KPI_Fastställd, data = final_df %>% filter(kön != "totalt"))
summary(naive_model)
```

Call:

```
lm(formula = antal ~ år + kön + examen + medelinkomst + Folkmängd +
  AVG_KPI_Fastställd, data = final_df %>% filter(kön != "totalt"))
```

Residuals:

Min	1Q	Median	3Q	Max
-7336.1	-2820.9	-343.8	2509.4	20567.5

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.674e+06	6.974e+05	-2.400	0.01733 *
år	8.648e+02	3.656e+02	2.365	0.01898 *
könmän	-1.241e+03	3.123e+03	-0.397	0.69157
examenLicentiatexamen	-2.188e+04	8.268e+02	-26.464	< 2e-16 ***
examenDoktorsexamen	-2.129e+04	8.268e+02	-25.744	< 2e-16 ***
medelinkomst	-3.498e+01	3.875e+01	-0.903	0.36787
Folkmängd	-1.050e-03	8.651e-03	-0.121	0.90357
AVG_KPI_Fastställd	-7.788e+01	2.982e+01	-2.611	0.00972 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4821 on 196 degrees of freedom

Multiple R-squared: 0.836, Adjusted R-squared: 0.8301

F-statistic: 142.7 on 7 and 196 DF, p-value: < 2.2e-16

```
[152]: naive_model2 <- lm(antal ~ år + kön + examen + medelinkomst +
  ↪AVG_KPI_Fastställd, data = final_df %>% filter(kön != "totalt"))
summary(naive_model2)
```

Call:

```
lm(formula = antal ~ år + kön + examen + medelinkomst + AVG_KPI_Fastställd,
  data = final_df %>% filter(kön != "totalt"))
```

Residuals:

Min	1Q	Median	3Q	Max
-7296	-2807	-362	2527	20515

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-1.628e+06	5.878e+05	-2.770	0.00613 **
år	8.397e+02	3.006e+02	2.793	0.00573 **

```

könmän                -1.184e+03  3.080e+03  -0.384  0.70108
examenLicentiatexamen -2.188e+04  8.247e+02 -26.531  < 2e-16 ***
examenDoktorsexamen   -2.129e+04  8.247e+02 -25.809  < 2e-16 ***
medelinkomst           -3.458e+01  3.852e+01  -0.898  0.37042
AVG_KPI_Fastställd     -7.729e+01  2.934e+01  -2.634  0.00911 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4809 on 197 degrees of freedom
Multiple R-squared:  0.8359,    Adjusted R-squared:  0.8309
F-statistic: 167.3 on 6 and 197 DF,  p-value: < 2.2e-16

```

This naive model results in a relatively high R^2 and R^2_{adj} values of $R^2 = 0.836$, $R^2_{adj} = 0.830$. The most significant regressors were the average KPI that year and if the degree was a “Licenciat” or PH. That the type of degree was significant is not at all unexpected as a Bachelor’s level degree is a prerequisite for further studies. That KPI was significant could hint at the state of the economy being important in people studying or not. But the model is quite hard to interpret as a person, because of the large amount of explanatory variables. Because of this the regressor with the highest p -value was removed from the model, this being the population variable.

The new model had $R^2 = 0.836$, $R^2_{adj} = 0.831$ meaning it performed slightly better than the naive model. Although it should be noted that all of the p -values barely changed.

In order to increase the interpretability of the model and to more easily choose our variables, we will instead create two new models one for the amount of people graduating with a Bachelor’s level degree and another for the amount of people graduating with either a “Licenciat” or PHD.

We will also approach the problem of choosing variables in the opposite way from what was done to the naive model, that is we will add one variable at a time instead of removing them.

```

[153]: bachelor_model <- lm(antal ~ ., data = final_df %>% filter(examen == "Grundutbildning", kön != "totalt") %>% select(!examen))
      ols_step_forward_p(bachelor_model)

higher_ed_model <- lm(antal ~ ., data = final_df %>% filter(examen != "Grundutbildning", kön != "totalt") %>% mutate(examen = factor(examen)))
      ols_step_forward_p(higher_ed_model)

```

Stepwise Summary						
Step	Variable	AIC	SBC	SBIC	R2	Adj.
↪ R2						
0	Base Model	1447.805	1452.244	1251.378	0.00000	0.
↪ 00000						
1	Folkmängd	1388.922	1395.580	1191.690	0.59153	0.
↪ 58534						

2	kön	1339.183	1348.061	1142.668	0.80914	0.
↩80326						
3	medelinkomst	1327.645	1338.743	1131.189	0.84359	0.
↩83626						
4	AVG_KPI_Fastställd	1320.904	1334.221	1124.602	0.86246	0.
↩85373						
5	år	1291.093	1306.629	1099.260	0.91385	0.
↩90690						

Final Model Output

Model Summary

R	0.956	RMSE	2897.053
R-Squared	0.914	MSE	8392914.012
Adj. R-Squared	0.907	Coef. Var	13.662
Pred R-Squared	0.896	AIC	1291.093
MAE	2360.274	SBC	1306.629

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error

AIC: Akaike Information Criteria

SBC: Schwarz Bayesian Criteria

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	6053846173.952	5	1210769234.790	131.532	0.0000
Residual	570718152.798	62	9205131.497		
Total	6624564326.750	67			

Parameter Estimates

	model	Beta	Std. Error	Std. Beta	t	Sig.
↩	lower	upper				↩
	(Intercept)	-4744231.247	760146.820		-6.241	0.000↩
↩	-6263743.089	-3224719.405				
	Folkmängd	0.001	0.009	0.021	0.113	0.911↩
↩	-0.018	0.020				

	könmän	-4492.417	3404.323	-0.228	-1.320	0.192
↪	-11297.563	2312.728				
	medelinkomst	-98.245	42.241	-0.536	-2.326	0.023
↪	-182.683	-13.807				
AVG_KPI_Fastställd		-231.133	32.509	-1.486	-7.110	0.000
↪	-296.118	-166.148				
	år	2423.482	398.513	2.409	6.081	0.000
↪	1626.866	3220.097				

Stepwise Summary

Step	Variable	AIC	SBC	SBIC	R2	Adj.
↪ R2						
0	Base Model	2054.785	2060.611	1665.554	0.00000	0.
↪ 00000						
1	examen	1980.883	1989.621	1590.585	0.42771	0.
↪ 42344						
2	medelinkomst	1863.077	1874.728	1473.845	0.76284	0.
↪ 75928						
3	AVG_KPI_Fastställd	1817.731	1832.294	1429.629	0.83257	0.
↪ 82876						
4	Folkmängd	1804.162	1821.638	1416.537	0.85068	0.
↪ 84612						
5	år	1784.163	1804.551	1398.030	0.87298	0.
↪ 86809						
6	kön	1776.912	1800.214	1391.715	0.88133	0.
↪ 87581						

Final Model Output

Model Summary

R	0.939	RMSE	156.805
R-Squared	0.881	MSE	24587.845
Adj. R-Squared	0.876	Coef. Var	25.765
Pred R-Squared	0.868	AIC	1776.912
MAE	132.207	SBC	1800.214

RMSE: Root Mean Square Error

MSE: Mean Square Error

MAE: Mean Absolute Error
AIC: Akaike Information Criteria
SBC: Schwarz Bayesian Criteria

ANOVA

	Sum of Squares	DF	Mean Square	F	Sig.
Regression	24834911.666	6	4139151.944	159.677	0.0000
Residual	3343946.892	129	25922.069		
Total	28178858.559	135			

Parameter Estimates

	model	Beta	Std. Error	Std. Beta	t	Sig.
	lower	upper				
	(Intercept)	-160197.705	28523.471		-5.616	0.000
	-216632.090	-103763.320				
examenDoktorsexamen		595.382	27.612	0.654	21.563	0.000
	540.752	650.013				
medelinkomst		-3.341	1.585	-0.395	-2.108	0.037
	-6.477	-0.205				
AVG_KPI_Fastställd		-1.257	1.220	-0.175	-1.030	0.305
	-3.670	1.157				
Folkmängd		-0.002	0.000	-0.905	-5.951	0.000
	-0.003	-0.001				
år		85.460	14.954	1.842	5.715	0.000
	55.874	115.046				
könmän		384.918	127.743	0.423	3.013	0.003
	132.176	637.660				

The variable selection algorithm for the Bachelor's level model resulted in the following choices: 1. Population 2. Gender 3. Average income 4. Average KPI 5. Year

On first glance this result was surprising to me as the first variable it chose was population in contrast to the naive model where it had the largest p -values, and that the variable with the lowest p -value except type of degree was added second to last.

Although from the pairs plot we knew that there was collinearity between some of the variables so it is to be expected that they would interact in nonintuitive ways.

For the model of the other two degrees, forward selection resulted in the choices: